

Malaria prevalence and child mortality — Methods & Results

Malaria prevalence and child mortality — Methods & Results summary

Three analyses relating *P. falciparum* prevalence to all-cause under-5 mortality in sub-Saharan Africa. All figures/tables below are reproduced by `run_all.R` into `results/`.

Data sources: MAP PfPR₂₋₁₀ raster (2024); IHME/GBD under-5 malaria deaths (2025); UN IGME all-cause child mortality + World Bank live births, **GDP per capita** (NY.GDP.PCAP.CD) and **DTP3** (SH.IMM.IDPT); DHS/MIS microdata (PR + birth-history recodes) across 23 SSA countries.

Component 1 — Country level

Methods. Malaria’s share of all-cause under-5 deaths = IHME under-5 malaria deaths ÷ IGME all-cause under-5 deaths (U5MR/1000 × live births), computed for **all under-5** and **neonatal-excluded** (denominator = 1mo–5y = U5MR - neonatal). Exposure = population-weighted national **PfPR₂₋₁₀** (MAP 2024). Linear model `share ~ PfPR + log(GDP p.c.) + DTP3`; outlier countries flagged by studentized residual from the prevalence-only fit.

Results (42 SSA countries). Malaria’s share of child deaths rises strongly with prevalence: **≈ +0.9 percentage-points of the all-U5 share (≈ +1.3 pp of the 1mo–5y share) per +1 PfPR₂₋₁₀ point** (adjusted for GDP and DTP3; adjusted R² ≈ 0.5–0.6). **Burundi** is a striking outlier (studentized residual ≈ 8) — malaria a far higher share of child deaths than prevalence predicts — followed by Uganda, Rwanda and Ghana. GDP per capita is not significant once DTP3 is included.

Tables: `results/component1_model_coefficients.csv`, `component1_outliers.csv`.




Figure 1: Component 1: malaria share of child deaths vs national PfPR

Component 2 — DHS subnational + national

Methods. One observation per **survey-region** across all usable DHS/MIS surveys (malaria biomarker + birth history). All-cause **U5MR** and **1mo-5y** mortality from the birth histories via `DHS.rates::chmort`. Exposure = child malaria prevalence put on a common footing: **microscopy-equivalent** (RDT to microscopy via the Component-3 factor where microscopy is missing), then **age-standardised 0.5-5y to PfPR₂₋₁₀ (2-10y)** with `malariaAtlas::convertPrevalence` (Smith et al. 2007 age-prevalence model). Two specifications:

1. **Primary — log-rate linear mixed model:** $\log(\text{mortality}) \sim \text{PfPR}_{2-10} + \text{DTP3} + \log(\text{GDP}) + \% \text{urban} + \text{year} + (1 + \text{PfPR}_{2-10} || \text{country})$.
2. **Count — negative-binomial GAM (mgcv):** region under-5 deaths with a $\log(\text{exposure})$ offset, a smooth $s(\text{year}_c)$, and country random intercept + slope.

Child stunting was evaluated as a further confounder but **excluded from the primary model**: five MIS-type surveys (including Liberia's only survey) collected no anthropometry, so adding it dropped ~11% of survey-regions and one country while leaving the malaria coefficient essentially unchanged.

Results (600 survey-regions, 44 surveys, 23 countries). Higher child malaria prevalence predicts higher child mortality, adjusted for all covariates:

	LMM (log-rate)	GAM (nb count)
U5MR , per +10	+7.0% (95% CI +3.2,	+6.9% (+3.5, +10.4)
PfPR ₂₋₁₀ pts	+10.9)	

	LMM (log-rate)	GAM (nb count)
1mo–5y, per +10 pts	+11.3% (+5.8, +17.0)	+11.0% (+5.8, +16.5)

The effect is **positive in every country** (country random slopes; figure below) and covariates behave as expected — DTP3 and % urban protective, a $\sim 3\%/yr$ secular decline, GDP $\sim$ null after the others. **Sensitivity** restricting to mid-transmission regions (PfPR₂₋₁₀ 5–50%, 354 regions) attenuates the effect (LMM: U5MR +6.2%, 1mo–5y +8.4%; GAM: +5.6% / +8.7%) — positive but weaker/less precise over the narrower range.




Figure 2: Component 2: adjusted country-specific prevalence–mortality slopes

Exposure-response shape. Replacing linear prevalence with a low-df smooth $s(\text{PfPR}_{2-10}, k=4)$: **U5MR is linear** across the range (edf = 1.0 — a constant proportional effect), while **1mo–5y shows mild curvature** (edf ≈ 2.1 — a smooth, decelerating rise, steeper at low prevalence). Removing neonatal deaths from the denominator uncovers the modest nonlinearity the all-U5 relationship hides.

Tables: `results/component2_model_coefficients.csv`, `component2_count_model_coefficients.csv` (`sample = full / pfpr_5_50`), `component2_country_slopes.csv`, `component2_exposure_response_edf.csv`.
Sensitivity figure: `component2_sensitivity_5to50.png`.

Negative control — neonatal mortality. Malaria kills post-neonatal children, not neonates, so a genuine malaria signal should be **absent** in the neonatal window; a null there also argues the prevalence–mortality gradient is not just




Figure 3: Component 2: smooth exposure-response $s(\text{PfPR}_{2-10})$

general socio-economic confounding. Refitting the same model by age window gives **all under-5 +7.0%** and **1mo-5y +11.3%**, but **neonatal -0.1% (95% CI -3.2 to +3.1, $p = 0.96$)** per +10 PfPR_{2-10} points (GAM: +0.3%, $p = 0.81$). The effect is age-specific in the malaria-plausible direction, which strengthens the causal reading.




Figure 4: Component 2 negative control: prevalence effect by age window

At the **raw-data** level the crude neonatal-prevalence trend is *weakly positive* — but this is between-country confounding (poorer, higher-mortality countries have both higher prevalence and higher neonatal mortality); the **covariate- and country-adjusted** effect is null. So adjustment removes the (confounding-driven) neonatal signal while leaving the post-neonatal (malaria) signal intact:

Tables: `results/component2_negative_control.csv`; figure `results/component2_neonatal_scatter.png`.




Figure 5: Neonatal mortality vs prevalence, by survey-region

Component 3 — RDT vs microscopy (supporting)

Methods. Across all DHS/MIS survey-regions with both tests, regress microscopy prevalence on RDT prevalence to a consensus conversion factor, used by Component 2 to impute microscopy where only RDT exists.

Results. RDTs over-detect (persistent HRP2 antigen), so microscopy runs lower: $\text{microscopy} \approx 0.74 \times \text{RDT}$ (Pearson $r = 0.90$, 431 survey-regions).

Table: `results/rdt_microscopy_conversion.csv`.

Cross-component prediction — total child mortality, 10% to 30% PfPR₂₋₁₀

Holding covariates at their means and anchoring both models to the same baseline at 10% PfPR (within each outcome), the predicted change in **total** all-cause child mortality as prevalence rises from 10% to 30%:

Outcome	baseline @10%	Component 1	Component 2
		(share, non-malaria fixed)	(direct model)
All under-5 (U5MR)	64.3 / 1,000	+ 23.8% to 79.6	+ 14.4% to 73.6

Outcome	baseline @10%	Component 1	Component 2
		(share, non-malaria fixed)	(direct model)
1mo–5y (neonatal excl.)	38.3 / 1,000	+ 48.3% to 56.8	+ 23.8% to 47.4

Component 2 (direct total-mortality model) gives the empirical net effect; Component 1 (malaria’s share, assuming non-malaria mortality is fixed) gives the upper “malaria-as-pure-addition” bound. Effects are larger for 1mo–5y because the malaria-irrelevant neonatal floor is removed. The 95% CIs (shaded) overlap, so the two independent routes are consistent. (*1mo–5y is nonlinear — see the exposure-response above — so its log-linear prediction is an approximation over this span.*)

Malaria-attributable fraction of child deaths

From the Component 2 log-rate models, the proportion of child deaths attributable to malaria at prevalence p is $\mathbf{AF}(p) = 1 - \exp(-(\eta(p) - \eta(0)))$ — the fraction that would be averted if transmission fell to zero, covariates fixed. Because the model is multiplicative, **this fraction depends only on p , not on the mortality level N** ; malaria-attributable deaths per 1,000 = $N \times \mathbf{AF}(p)$. Shown for both outcomes as the **linear** closed form (LMM coefficient) and the **spline** (low-df nb-GAM $\mathbf{s}(\text{PfPR}, \mathbf{k}=4)$); the neonatal-excluded denominator uses $N = \text{U5MR} - \text{NNMR}$.

PfPR ₂₋₁₀	All under-5, linear	All under-5, spline	1mo–5y, linear	1mo–5y, spline
10%	6.5%	6.5%	10.1%	14.6%
20%	12.6%	12.5%	19.2%	25.7%
30%	18.3%	18.2%	27.4%	33.3%
40%	23.6%	23.5%	34.7%	38.4%
50%	28.6%	28.4%	41.3%	42.2%

For **U5MR** the linear and spline curves coincide (exposure-response is linear, edf ≈ 1.0). For **1mo–5y** the smooth (edf ≈ 2.1) rises a little faster than the linear form at low–mid prevalence then flattens, the two converging near ~40% by 50% PfPR — a mild, decelerating nonlinearity rather than the over-flexible wiggle a high-df spline produced. Removing neonatal deaths (largely non-malarial) still roughly doubles malaria’s share versus all-U5. These are




Figure 6: Component 3: RDT vs microscopy




Figure 7: Predicted total child mortality vs prevalence, 10% to 30% PfPR2-10




Figure 8: Malaria-attributable fraction of child deaths vs PfPR2-10

counterfactual (vs zero-transmission) fractions assuming the adjusted association is causal; the CIs (shaded) reflect only the prevalence-effect sampling error. Table: `results/attributable_fraction.csv`.

Comparing the two methods

Both components estimate the same quantity — malaria as a **% of child deaths** — so they can be overlaid on shared axes. Component 1 is the country-level share (points = $IHME \div IGME$ per country; line = fitted $share \sim PfPR + \log GDP + DTP3$ at mean covariates); Component 2 is the DHS attributable-fraction spline. Drawn over Component 1’s national-prevalence range (to ~40%).

The routes **broadly agree** and their 95% bands overlap throughout. For **all under-5** the two track the country cloud closely, Component 2 running slightly below Component 1. For **1mo–5y** both rise more steeply once neonatal deaths are removed; they coincide at low–mid transmission and diverge at the top of the range, where Component 1 extrapolates roughly linearly while Component 2 begins to plateau. That two independent data sources (IHME/IGME country totals vs DHS survey mortality) and two different models land on a similar prevalence–share relationship is the main cross-check of the work.

National malaria-burden estimates ($PfPR_{2-10} > 10\%$)

Applying both methods to a common denominator — all-cause child deaths = mortality rate $\times$ live births (IGME/WB) — gives the number of malaria-




Figure 9: Method comparison: Component 1 share vs Component 2 attributable fraction

attributable child deaths per country. **Component 1** multiplies by the fitted malaria *share* ($\text{share} \sim \text{PfPR} + \log \text{GDP} + \text{DTP3}$); **Component 2** by the attributable fraction $\text{AF} = 1 - \exp(-\beta \cdot \text{PfPR}/10)$ from the DHS log-rate model. Both are shown for **all under-5** and **1mo–5y (neonatal-excluded)**, with the outcome-matched coefficients. Only the malaria coefficient carries uncertainty (prevalence, GDP, DTP3, births, rates fixed); because that coefficient is shared across countries, a pooled-total CI is the total re-evaluated at the coefficient's bounds. A national/subnational check for Nigeria and DR Congo (below) showed aggregating admin-1 estimates moves the total by <1%, so national PfPR is the representative scale.

Malaria-attributable child deaths for the **25 countries** with national $\text{PfPR}_{2-10} > 10\%$ (point estimates; sorted by Component 2 all-U5). C1 = fitted share, C2 = attributable fraction:

Country	PfPR	C1 (U5)	C2 (U5)	C1 (1mo– 5y)	C2 (1mo– 5y)	IHME
Nigeria	24.7%	178,352	133,962	189,730	134,093	150,493
DR Congo	36.0%	125,535	86,371	146,432	94,214	68,300
Niger	17.3%	26,496	13,657	31,259	14,716	23,329
Uganda	21.9%	21,246	11,502	19,985	9,826	28,474
Mozambique	23.9%	17,193	11,324	16,013	9,892	11,388
Cameroon	25.2%	13,727	9,807	13,872	9,137	17,718
Angola	20.4%	8,394	8,841	8,550	8,489	15,266

Country	PfPR	C1 (U5)	C2 (U5)	C1 (1mo– 5y)	C2 (1mo– 5y)	IHME
Côte d’Ivoire	21.6%	11,184	8,806	10,665	7,695	14,113
Chad	15.0%	10,901	8,232	12,325	8,704	5,498
Mali	17.8%	13,424	7,988	13,846	7,533	15,377
Benin	36.4%	10,138	7,854	10,176	7,388	8,123
Burkina Faso	20.9%	13,717	7,254	15,672	7,529	12,439
Guinea	21.4%	6,608	6,077	7,160	6,247	7,979
South Sudan	23.1%	7,040	4,861	6,778	4,357	4,282
Central African Rep.	33.7%	5,179	4,487	5,230	4,446	4,320
Sierra Leone	30.8%	8,044	4,414	9,076	4,533	7,181
Malawi	19.2%	8,494	3,991	7,334	3,145	6,316
Ghana	16.2%	5,914	3,305	5,074	2,515	7,254
Burundi	23.4%	7,093	3,207	6,937	2,887	15,659
Zambia	14.3%	6,196	3,091	5,934	2,655	5,272
Togo	21.0%	3,570	1,993	3,575	1,810	1,972
Liberia	15.1%	2,656	1,432	2,774	1,363	3,088
Congo-Brazzaville	23.0%	1,439	1,083	1,374	948	1,414
Equatorial Guinea	22.8%	456	549	458	511	724
Gabon	18.4%	118	262	96	200	262
TOTAL (25)	—	513,116	354,352	550,323	354,836	436,242

Pooled-total 95% CIs (coefficient-only): all-U5 — C1 [357,649–668,584], C2 [172,582–519,653]; 1mo–5y — C1 [367,934–732,712], C2 [200,714–489,777]. All-cause denominators (2,292,071 U5; 1,507,890 post-neonatal) and every cell’s CI are in `results/country_malaria_deaths_gt10.csv`.

The IHME reference (436k, all-U5) sits between the two model routes (C1 ~18% above, C2 ~19% below). Component 2 is internally consistent across denominators (355k either way — the higher post-neonatal fraction offsets the smaller base); Component 1’s two share models are fit separately and differ by ~7%. Nigeria + DR Congo alone are ~60% of the burden. Per-country estimates (all columns, with CIs) are in `results/country_malaria_deaths_gt10.csv`; the largest single-country divergences from IHME are **Burundi** (IHME far higher than either model) and **Chad** (lower).

National vs subnational (admin-1) check — all under-5:

	National	Subnational (Σ admin-1)
Nigeria — Component 1	178,352	178,350
Nigeria — Component 2	133,962	133,488
DR Congo — Component 1	125,535	125,535
DR Congo — Component 2	86,371	85,696

Component 1 is unchanged (its share model is linear in prevalence); Component 2 is marginally lower subnationally (the AF curve is concave, so disaggregating high/low-prevalence units loses a little via Jensen’s inequality) — but within-country prevalence heterogeneity is small, so the effect is <1%. Scripts: R/07_ng_drc_malaria_deaths.R (results/ng_drc_malaria_deaths.csv), R/08_country_malaria_deaths.R (results/country_malaria_deaths_gt10.csv).

RCT triangulation — does the relationship hold in randomised trials?

As an external check, we asked whether the model-implied prevalence to mortality relationship holds in **randomised** insecticide-treated-net (ITN) trials. Using the Pryce et al. (2018) Cochrane ITN review, we took every sub-Saharan **cluster-RCT** that measured **both** a change in child parasite prevalence and all-cause under-5 mortality, extracting the numbers from the **original publications**: **D’Alessandro 1995** (Gambia, 5 zones), **Habluetzel 1997/99** (Burkina Faso), **Phillips-Howard 2003 / ter Kuile** (western Kenya) and **Nevill 1996 / Snow** (Kilifi). Binka 1996 measured no parasite prevalence (excluded); Snow 1987 is individually randomised (excluded).

Method. Each arm’s parasite prevalence is age-standardised to PfPR₂₋₁₀, and **Component 2** predicts the mortality reduction from the observed prevalence reduction by evaluating its **nonlinear spline at both the control and intervention prevalence levels**. (Component 1’s linear share model is unusable per-arm here — extrapolated to these prevalences it implies shares > 100%.)

Trial (zone)	Baseline PfPR ₂₋₁₀	Prevalence reduction (pts)	Observed U5 ↓	Component 2 predicted
D’Alessandro z1	37%	9	62%	7%
D’Alessandro z2	34%	10	51%	10%

Trial (zone)	Baseline PfPR ₂₋₁₀	Prevalence reduction (pts)	Observed U5 ↓	Component 2 predicted
D'Alessandro z3	26%	10	51%	11%
D'Alessandro z4	54%	10	5%	6%
D'Alessandro z5	46%	-27 (rose)	-20%	-16%
Habluetzel	95%	8	15%	4%
Phillips- Howard	77%	14	15%	7%
Nevill	39%	20	30%	18%




Figure 10: RCT triangulation: dose-response and Component 2 predicted vs observed

Findings. The trials show the expected **positive dose-response** — bigger prevalence reductions go with bigger mortality reductions — at an inverse-variance-weighted **≈ 8.6% fewer U5 deaths per 10 PfPR₂₋₁₀ points removed**, close to Component 2's own **≈ 11.3% per 10 points**. Component 2 correctly predicts the **negative control** (D'Alessandro zone 5, where the programme failed: prevalence *rose* and mortality *rose*) and the low-effect points (zone 4, Habluetzel, Phillips-Howard). It **under-predicts** the largest observed reductions (D'Alessandro zones 1–3, Nevill), but those rest on few deaths, infant-age or near-saturated prevalence, and their **95% CIs overlap** the model line. The experimental evidence is thus broadly **consistent** with the observational Component 2 relationship. Exploratory (8 heterogeneous points: curtains, an untreated-net comparator, spillover, infant-age; D'Alessandro-zone CIs are Poisson, un-clustered). Data: `results/triangulation_rct_data.csv`;

script R/10_triangulation_rct.R.

Caveats

- **Component 1 is ecological** (country-level); **Component 2** adjusts for the main confounders but is observational, and region-level U5MR from a single survey is noisy.
- MAP PfPR₂₋₁₀ is 2024 vs the 2025 IHME export (1-year offset).
- IHME (GBD) and WHO differ substantially on malaria mortality; Component 1's numerator is IHME.
- DHS prevalence is age-standardised to PfPR₂₋₁₀ via a model (Smith 2007), and RDT is converted to a microscopy footing; both are documented approximations.